

University of Southern California

Viterbi School of Engineering

EE352

Computer Organization and Architecture

Exceptions

References:

- 1) Textbook
- 2) Mark Redekopp's slide series

Shahin Nazarian

Spring 2010

What are Exceptions?

- Any event that causes a break in normal execution
 - Error Conditions
 - Invalid address, Arithmetic/FP overflow/error
 - Hardware Interrupts / Events
 - Handling a keyboard press, mouse moving, USB data transfer, etc.
 - System Calls / Traps
 - User applications calling OS code

Error Condition Exceptions


- Bus Error
 - No memory or I/O device responds to a read or write
- Address Error
 - Unaligned address used for half or word accesses
 - Access to Kernel memory space when in User mode
- Floating Point Exceptions
 - Many possible exceptions encompassed by this
- Integer Overflow
 - An instruction that causes 2's complement overflow
- TRAP Instructions
 - SW definable error conditions
- Reserved Instruction
 - User code attempting to perform kernel mode operations
- Coprocessor Unusable
 - Use of a coprocessor instruction when the coprocessor is not implemented

I/O Notification

- Two methods for I/O devices to indicate they need the processor's attention
 - **Polling** "busy" loop (responsibility on proc.)
 - Processor has responsibility of checking each I/O device
 - Many I/O events happen infrequently (1-10 ms) with respect to the processors ability to execute instructions (1-10 ns)
 - **Interrupts** (responsibility on I/O device)
 - I/O device notifies processor only when it needs attention

```
Getkey: la      $t0,KEY_STAT ; get status
        lw      $t1,0($t0)
        andi    $t1,$t1,0x0001
        beq     $t1,$zero,getkey
        la      $t2,KEY_DATA ; get key
        lw      $t3,0($t2)
```

Polling Loop



Interrupt

System Calls / TRAP Exceptions

- Provide a controlled method for user mode applications to call kernel mode (OS) code
- **Syscalls** and **traps** are very similar to subroutine calls but they switch into **kernel mode** when called
 - Kernel mode is a special mode of the processor for executing trusted (OS) code

Exception Processing


- Save necessary state to be able to restart the process
 - Save PC of offending instruction
- Call an appropriate **handler** routine to deal with the error / interrupt / syscall
 - Handler identifies the cause of exception and handles it
 - May need to save more state
- Restore state and return to offending application (or kill it if recovery is impossible)

Coprocessor 0 Registers

- **Status Register**
 - Enables and disables the handling of exceptions/interrupts
 - Controls user/kernel processor modes
 - Kernel mode allows access to certain regions of the address space and execution of certain instructions
- **Cause Register:** Indicates which exception/interrupt occurred
- **Exception PC (EPC) Register**
 - Indicates the address of the instruction causing the exception
 - This is also the instruction we should return to after handling the exception
- Coprocessor registers can be accessed via the '**mtc0**' and '**mfc0**' instructions
 - `mfc0 $gpr,$c0_reg` `# R[gpr] = C0[c0_reg]`
 - `mtc0 $gpr,$c0_reg` `# C0[c0_reg] = R[gpr]`

Status Register


- Register 12 in coprocessor 0
- Bit definitions
 - **IM**[7:0] - Interrupt Mask
 - 1 = Ignore interrupt level / 0 = Allow interrupt level
 - **UM** - User Mode
 - 1 = User mode / 0 = Kernel mode
 - **ERL/EXL** = Exception/Error Level
 - 1 = Already handling exception or error / 0 = Normal exec.
 - If either bit is '1' processor is also said to be in kernel mode
 - **IE** = Interrupt Enable
 - 1 = Allow unmasked interrupts / 0 = Ignore all interrupts despite the IM bits



Cause Register


- Register 13 in coprocessor 0
- Bit definitions
 - **BD - Branch Delay**
 - The offending instruction was in the branch delay slot
 - EPC points at the branch but it was EPC+4 that caused the exception
 - **PI[7:0] - Pending Interrupt**
 - 1 = Interrupt Requested / 0 = No interrupt requested
 - **Exception Code** - Indicates cause of exception (see table)

Code	Cause
0	Interrupt (HW)
4, 5	Load (4), Store (5) Address Error
6, 7	Instruc. (6), Data (7) Bus Error
8	Syscall
9	Breakpoint
10	Reserved Instruc.
11	CoProc. Unusable
12	Arith. Overflow
13	Trap
15	Floating Point



EPC Register

- Exception PC holds the address of the offending instruction (unless BD bit is set)
 - Can be used along with 'Cause' register to find and correct some error conditions
- **'eret'** instruction used to return from exception handler, back to execution point in original code (unless handling the error means having the OS kill the process)
 - **'eret' Operation: $PC = EPC$**



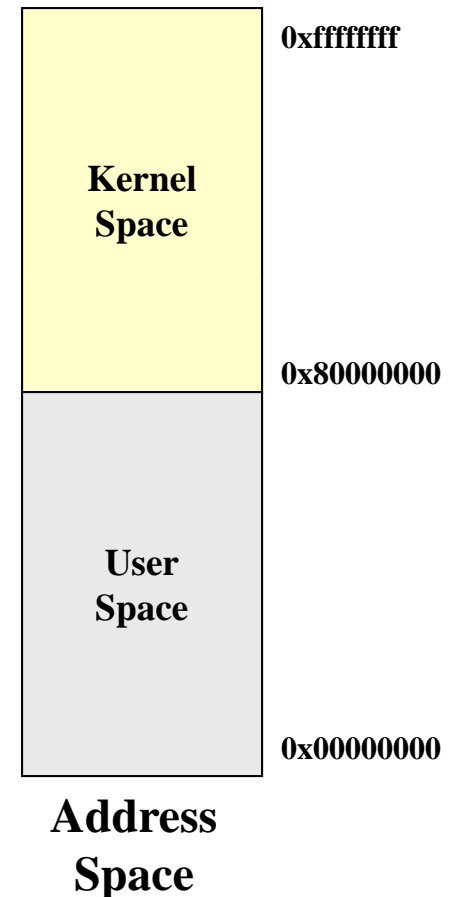
Address of instruction that generated the exception

Kernel and User Mode

- User applications are designed to run in user mode
- OS and other system software should run in kernel mode
- Certain features/privileges are only allowed to code running in kernel mode
- Processor is in kernel mode if
 $\{UM = 0 \text{ OR } ERL = 1 \text{ OR } EXL = 1\}$

Kernel Mode Privileges

- **Privileged instructions**
 - 'eret', 'di', 'ei', etc.
 - User apps. shouldn't be allowed to disable/enable interrupts, change memory mappings, etc.
- **Privileged Memory or I/O access**
 - Processor supports special areas of memory or I/O space that can only be accessed from kernel mode
- **Separate stacks and register sets**
 - MIPS processors can use "shadow" register sets (alternate GPRs when in kernel mode)



Syscalls


- Provides a structured entry point to the OS
 - Really just a subroutine call that also switches into kernel mode
 - Often used to allow user apps. to request I/O or other services from the OS
- Syntax: `syscall`
 - Necessary arguments are defined by the OS and expected to be placed in certain registers

Calling Exception Handlers

- Processor will automatically call a routine known as an **exception handler** to handle the exception
- Problems with calling exception handlers
 - How does the HW 'automatically' call handler routine since we don't know where we will be in our code when an exception occurs
 - Calling the exception handler can cause changes in state (registers) when we return to the running application

Problem of Calling a Handler

- We can't use explicit 'jal' instructions to call exception handlers since we don't know when they will occur

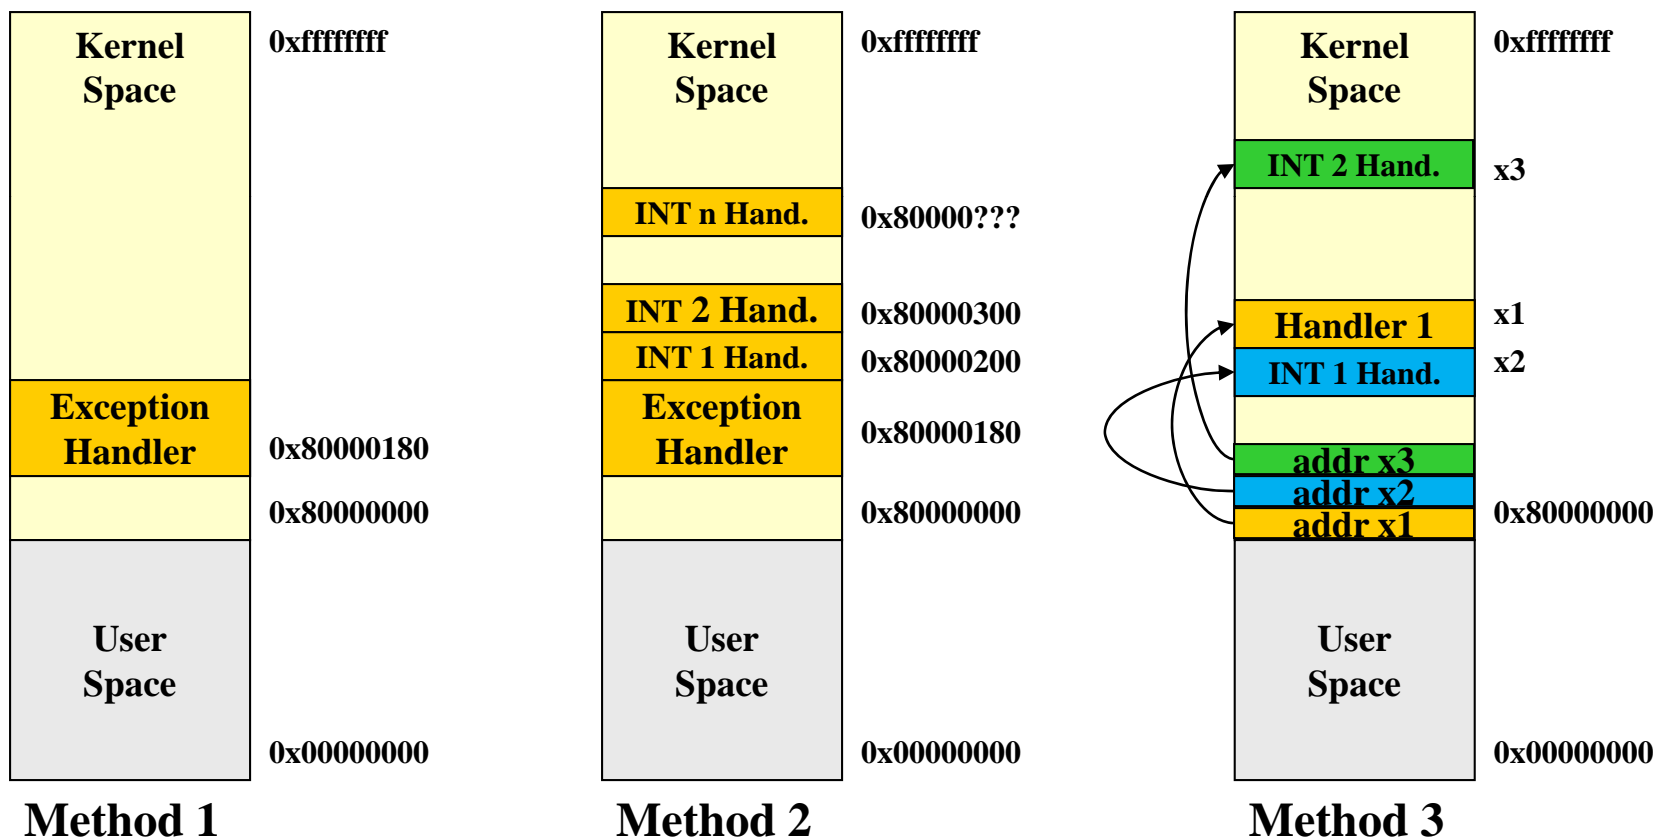


Many instructions could cause an error condition. Or a hardware event like a keyboard press could occur at any point in the code

Solution for Calling a Handler

- Since we don't know when an exception will occur there must be a preset location where an exception handler should be defined or some way of telling the processor in advance where our exception handlers will be located
- **Method 1: Single hardwired address for master handler**
 - Early MIPS architecture defines that the exception handler should be located at 0x8000_0180. Code there should then examine CAUSE register and then call appropriate handler routine (Mars uses this method)
- **Method 2: Vectored locations (usually for interrupts)**
 - Each interrupt handler at a different address based on interrupt number (a.k.a. vector) (INT1 @ 0x80000200, INT2 @ 0x80000300)
- **Method 3: Vector tables**
 - Table in memory holding start address of exception handlers (i.e. overflow exception handler pointer at 0x0004, FP exception handler pointer at 0x0008, etc.)

Handler Calling Methods



Exception Processing Steps

- Save PC (address of offending instruction) [HW]
 - So we know where to return once we're done handling the exception
- Record the cause of the exception [HW]
 - HW will identify source of HW errors or interrupts
- Switch to kernel mode [HW]
- Calculate and jump to exception handler [HW]
 - Using appropriate method for locating handler
- Execute Handler [SW]
 - Likely need to save/restore state (i.e. GPRs)
- Use 'eret' to return to address in EPC [SW]

Sample Exception Handler

- Main handler needs to examine cause register and call a more specific handler
- Handlers must end with 'eret' rather than 'jr \$ra'

```
.text
L1: li    $t0,0x100A1233
    lw    $s0,0($t0)
    ---
```

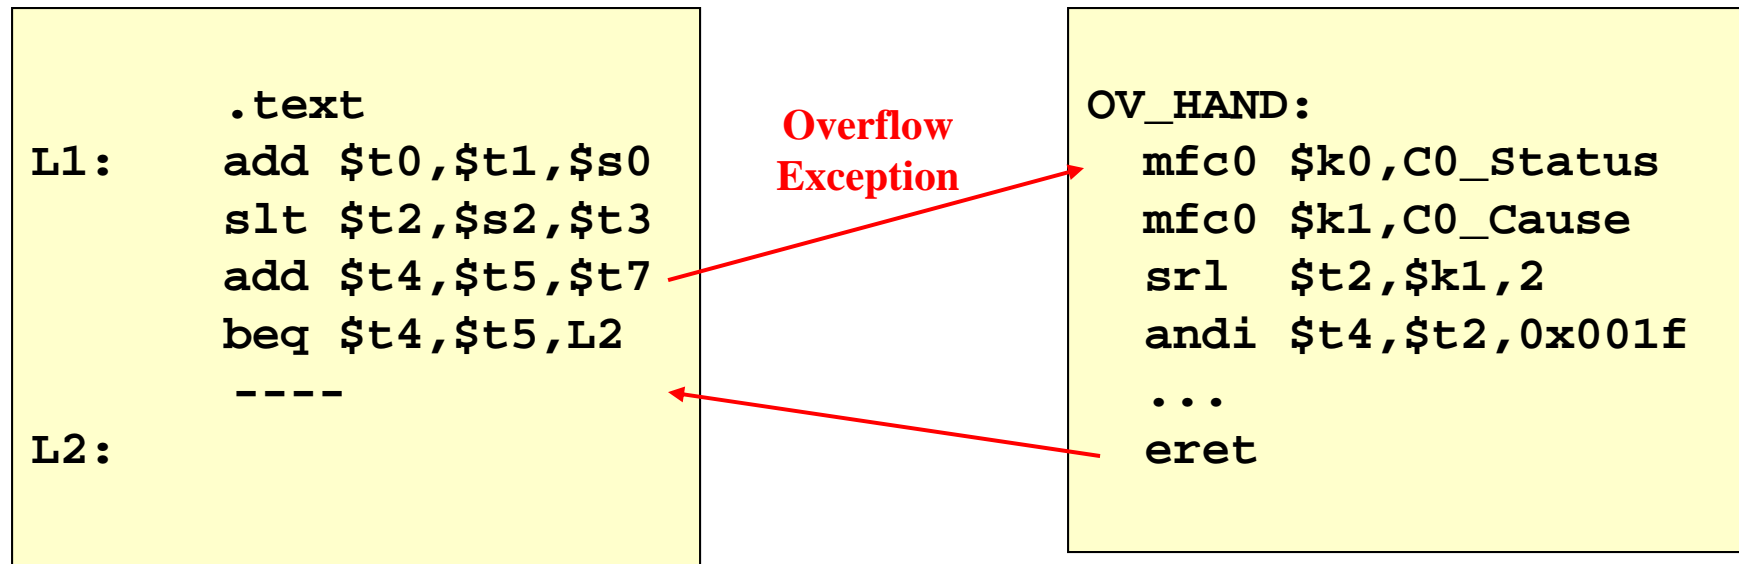
**Invalid Address
Exception**

```
0x8000_0180:
    mfc0   $k0,C0_Status
    mfc0   $k1,C0_Cause
    srl    $t2,$k1,2
    andi   $t4,$t2,0x001f
    bne    $t4,0,E1
    j      INT_HAND
E1: ...
E4: bne    $t4,4,E2
    j      ADDR_HAND
...
ADDR_HAND:
    ...
    eret
```

Main handler can determine cause and

Problem of Changed State

- Since exceptions can occur at any time, the handler **must save all registers it uses** (except for kernel registers \$k0-\$k1 which should be unused by the user app)



Handlers need to save/restore values to stack to avoid overwriting needed register values (e.g. \$t2, \$t4)

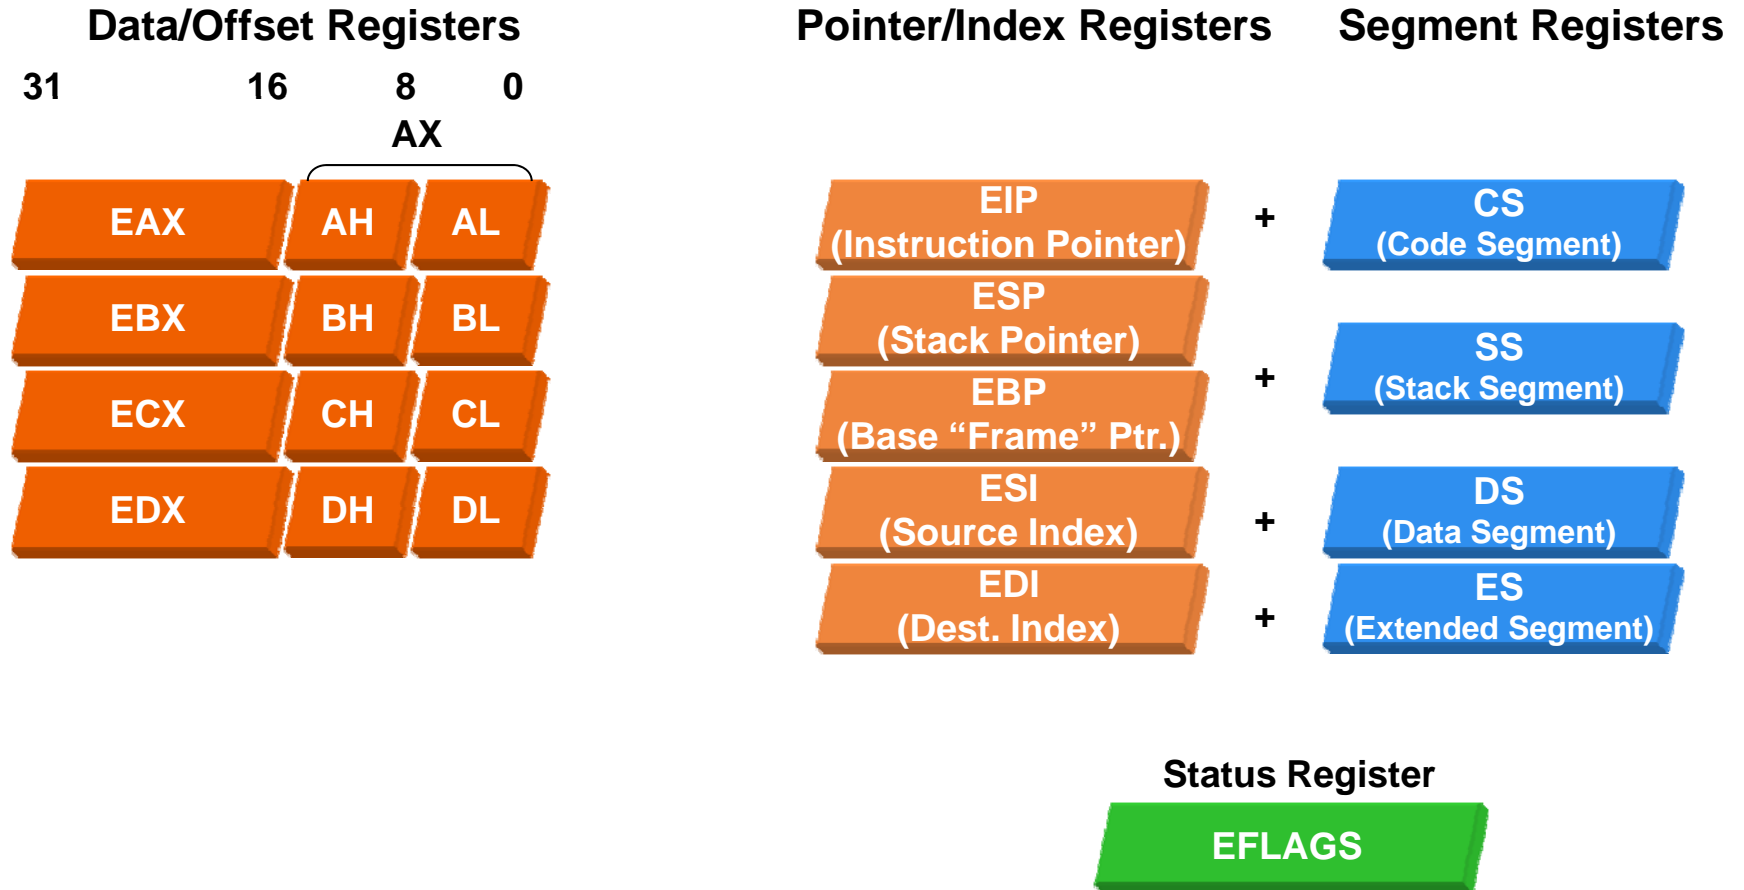
Bonus Material

Intel Architectures

Intel Architectures


Processor	Year	Address Size	Data Size
8086	1978	20	16
80286	1982	24	16
80386/486	'85/'89	32	32
Pentium	1993	32	32
Pentium 4	2000	32	32
Core 2 Duo	2006	64	64

Intel (IA-32) Architectures



Real Mode Addressing

- How to make 20-bit address w/ 16-bit registers?
 - Use 2 16-bit registers
 - (Segment register * 16) + Index Reg.
- Format:
 - **Seg Reg:Index Reg (e.g. CS:IP)**



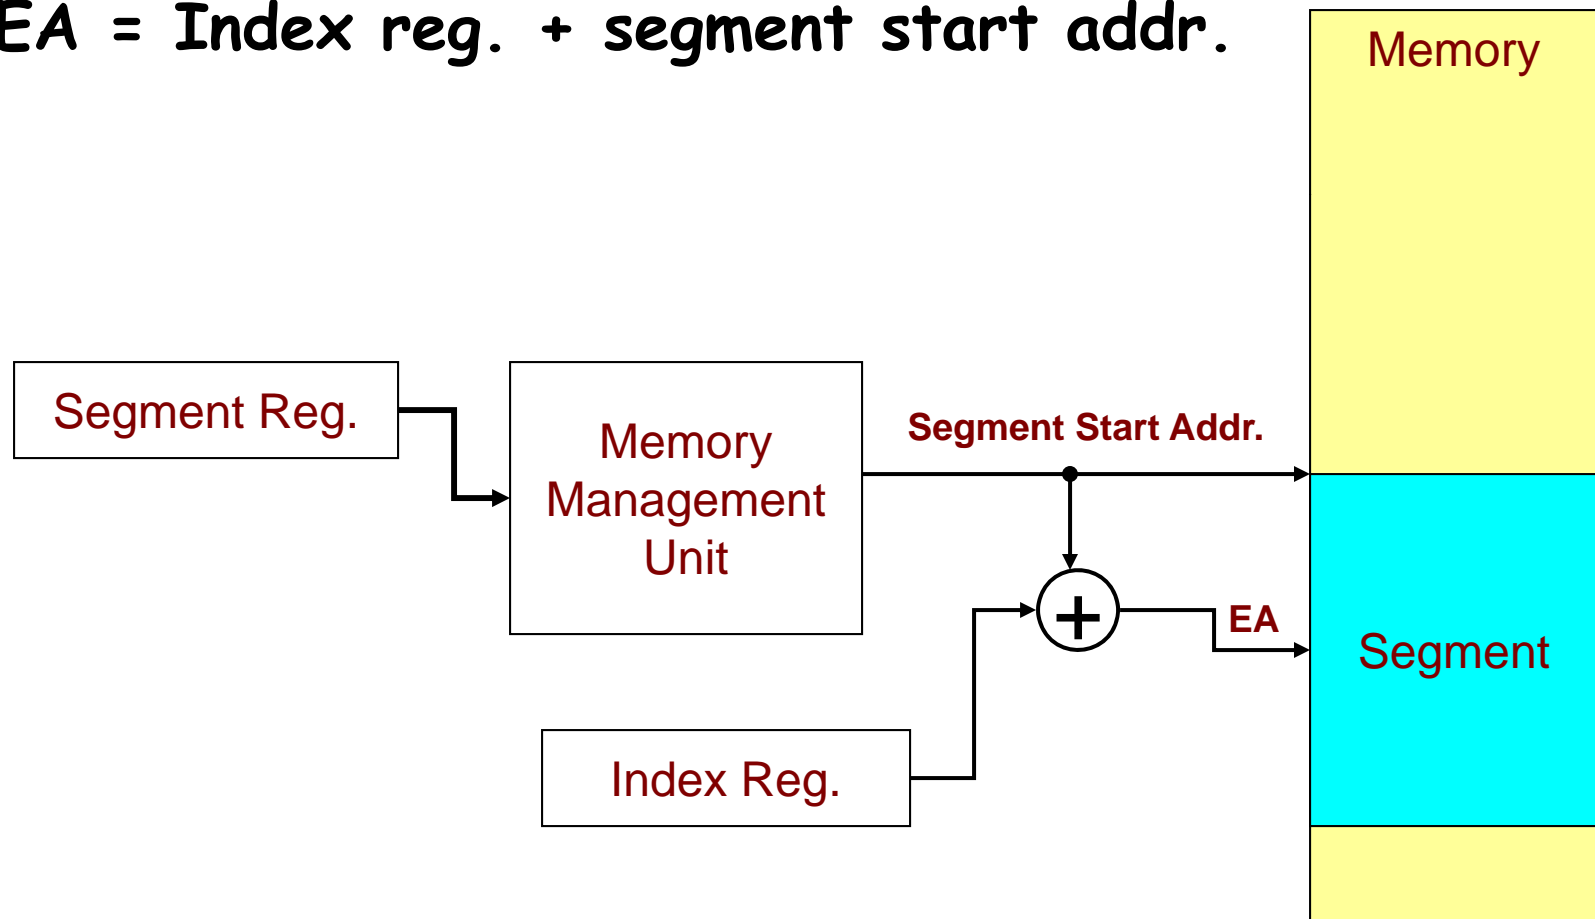
Examples:

$$\begin{array}{rcl} \$1A5C:\$405E & \rightarrow & \begin{array}{r} \$1A5C0 \\ + \$405E \\ \hline \$1E61E \end{array} \end{array}$$

$$\begin{array}{rcl} \$74AB:\$E892 & \rightarrow & \begin{array}{r} \$74AB0 \\ + \$E892 \\ \hline \$83345 \end{array} \end{array}$$

Protected Mode Addressing

- Segment Register value selects a segment (area of memory)
- $EA = \text{Index reg.} + \text{segment start addr.}$



IA-32 Addressing Modes

Name	Example	Effective Address
Immediate	MOV EAX,5	Operand = Value
Direct	MOV EAX,[100]	EA = Addr
Register	MOV EAX,EDX	EA = Reg
Register indirect	MOV EAX,[EBX]	EA = (Reg)
Base w/ Disp.	MOV EAX,[EBP+60]	EA = (Reg) + Disp
Index w/ Disp.	MOV EAX,[ESI*4+10]	EA = (Reg)*S + Disp.
Base w/ Index	MOV EAX,[EBP + ESI*4]	EA = (Reg1)+(Reg2)*S
Base w/ Index & Disp.	MOV EAX,[EBP+ESI*4+100]	EA = (Reg1)+(Reg2)*S + Disp.

IA-32 Instructions

- **Stack aware instructions**
 - 'call' / 'ret' automatically push/pop return address onto stack
 - 'push' / 'pop' instructions for performing these operations
- **Memory / Register architecture**
 - One operand can be located in memory
 - `add eax, [ebp]` # adds reg. EAX to value pointed at by EBP
- **Specialized Instructions**
 - `xchg src1, src2` # exchanges/swaps operands
 - string copy and compare instructions