

University of Southern California

Viterbi School of Engineering

EE352

Computer Organization and Architecture

**Exploiting Thread Level Parallelism
and Data Level Parallelism**

References:

- 1) Textbook
- 2) Mark Redekopp's slide series

Shahin Nazarian

Spring 2010

Exploiting Parallelism

- With increasing transistor budgets of modern processors (i.e., can do more things at the same time) the question becomes how do we find enough *useful* tasks to increase performance, or, put another way, what are the most effective ways of exploiting parallelism!
- Many types of parallelism are available
 - **Instruction Level Parallelism (ILP)**: Overlapping instructions within a single process/thread of execution
 - **Thread Level Parallelism (TLP)**: Overlap execution of multiple processes / threads
 - **Data Level Parallelism (DLP)**: Overlap an operation (instruction) that is to be applied to multiple data values (usually in an array)
 - For (i=0; i < MAX; i++) { A[i] = A[i] + 5; }



Data Level Parallelism

SIMD/Vector Units

Introductory Example for SIMD

- An image is just a 2-D array of pixel values
- Pixel color represented as a number (between 0 = black, to 255 = white)




Image taken from the photo "Robin Jeffers at Ton House" (1927) by Edward Weston

Shahin Nazarian/EE352/Spring10

Graphics Operations


- **Brightness**
 - Each pixel value is increased/decreased by a constant amount
 - $P_{\text{new}} = P_{\text{old}} + B$
 - $B > 0$ = brighter
 - $B < 0$ = less bright
- **Contrast**
 - Each pixel value is multiplied by a constant amount
 - $P_{\text{new}} = C * P_{\text{old}}$
 - $C > 1$ = more contrast
 - $0 < C < 1$ = less contrast
- Same operations performed on all pixels




+ Contrast



- Brightness



Original



+ Brightness



- Contrast

Vector Operations

- Need to perform single operation (instruction) such as $+B$ or $*C$ on multiple data values
- General-purpose computers use instructions that perform an operation on a single (“scalar”) data value (i.e. `lw $t0, Pold` / `add $t0, $t0, 5` / `sw $t0, Pnew`)
- Modern processors include HW capability & instructions that perform the same operation on multiple data items (a.k.a. “vectors”) using a single instruction
 - HW adds “vector” registers which can hold 128- or 256-bit data
 - 128- or 256-bit data can be interpreted as a set of 8, 16, or 32-data items
 - Example instruction: `'vaddw $v0, $v0, 5'` where `vaddw` stands for Vector Add Word and `$v0` is a 128-bit chunk

Vector Operations (Cont.)



- Computers that perform the same operation/instruction on several data values are referred to as **SIMD (Single Instruction, Multiple Data) machines** or **Vector/Stream processing machines**

Vector Processing

- New instructions and WIDE registers that can be selected to operate on different size data
 - Treat data as 8-bit, 16-bit, 32-bit chunks


Vector Instructions

- **Vector load / store**
 - “LV” reads a 128-bit chunk from memory to 128-bit vector register
 - “SV” writes a 128-bit chunk to memory from 128-bit vector register
- **Vector operations**
 - “VADDB” = (16) 8-bit values
 - “VADDH” = (8) 16-bit values
 - “VADDW” = (4) 32-bit values



Vector Operations

- `vaddw $v0,$v1,$v2`



More Vector Instructions

```
for(i=MAX; i != 0; i--)  
    A[i] = A[i] + 5;
```

Original “scalar” code

```
// Loop unrolled 4 times  
for(i=MAX; i != 0; i=i-4){  
    A[i] = A[i] + 5;  
    A[i+1] = A[i+1] + 5;  
    A[i+2] = A[i+2] + 5;  
    A[i+3] = A[i+3] + 5;  
}
```

Unrolled
Scalar
Code

```
$s0 = pointer to A[i]  
$s1 = i = # of iterations  
  
L1: lw    $t1,0($s0)  
     addi $t1,$t1,5  
     sw    $t1,0($s0)  
     lw    $t1,4($s0)  
     addi $t1,$t1,5  
     sw    $t1,4($s0)  
     lw    $t1,8($s0)  
     addi $t1,$t1,5  
     sw    $t1,8($s0)  
     lw    $t1,12($s0)  
     addi $t1,$t1,5  
     sw    $t1,12($s0)  
     addi $s0,$s0,16  
     addi $s1,$s1,-4  
     bne  $s1,$zero,L1
```

Vectorized / SIMD Code
(Could unroll this if desired)

```
$s0 = pointer to A[i]  
$s1 = i = # of iterations  
  
     lvi    $v2,5  
L1: lv     $v1,0($s0)  
     vaddw  $v1,$v1,$v2  
     sv     $v1,0($s0)  
     addi   $s0,$s0,16  
     addi   $s1,$s1,-4  
     bne    $s1,$zero,L1
```

Vector Processing Examples

- Remember: vector processing machines are SIMD type
- **Pentium**
 - **MMX (MultiMedia eXtension)** - 64-bit registers with only integer vector operations
 - **SSE (Streaming SIMD Extensions,)** SSE2, SSE3 - 128-bit vectors with support for single- and double-precision FP ops
- **PowerPC**
 - **Altivec** instructions extensions
- **Cell Processor**
 - Each of the 8 **SPEs (Synergistic Processing Elements)** use 128-bit vector operations
 - No scalar (single) data units...instead just don't use n-1 slots of vector operations

Optional to Review

- **MMX**

[http://en.wikipedia.org/wiki/MMX_\(instruction_set\)](http://en.wikipedia.org/wiki/MMX_(instruction_set))

- **SSE**

http://en.wikipedia.org/wiki/Streaming_SIMD_Extensions

- **AltiVec**

<http://en.wikipedia.org/wiki/AltiVec>

- **Cell Processor**

- Check **SPE** and **PPE** in the following:

[http://en.wikipedia.org/wiki/Cell_\(microprocessor\)#Synergistic_Processing_Elements_.28SPE.29](http://en.wikipedia.org/wiki/Cell_(microprocessor)#Synergistic_Processing_Elements_.28SPE.29)




Thread Level Parallelism

**Mulithreaded and Multicore
Processors**

Limit to ILP

- ILP has limits due to actual data (RAW) dependencies
- Impractical amount of HW needed to get close to this limit
- VLIW (Very Large Instruction Word) processors take advantage of ILP



© 2007 Elsevier, Inc. All rights reserved.

Thread-Level Parallelism


- ILP attempts to find, fetch, and execute independent instructions from within a single thread
 - Requires a lot of HW resources and may not be possible
- TLP fetches instruction from different threads which are independent *by definition*
 - A thread: Separate logical stream of instructions + data
 - Separate instructions = Separate PC
 - Separate data = Separate register sets, stack, and possibly memory address space
- This way we can fill our issue slots w/o having to work so hard to find independent instructions

Multi-threaded Processors

- Requires duplicating and/or sharing HW resources to support multiple threads
- Duplicated HW
 - 2 PCs to be able to fetch instructions from different threads
 - Instructions will be “tagged” with `thread_id` to maintain separability
- Duplicated or shared HW
 - Logically separate registers (can be physically separate or shared)
 - `$t0` for thread A should be different from `$t0` for thread B
 - Branch Predictors, **TLBs*** (translation-lookaside buffer), etc.

Example: 2-way Threaded Machine

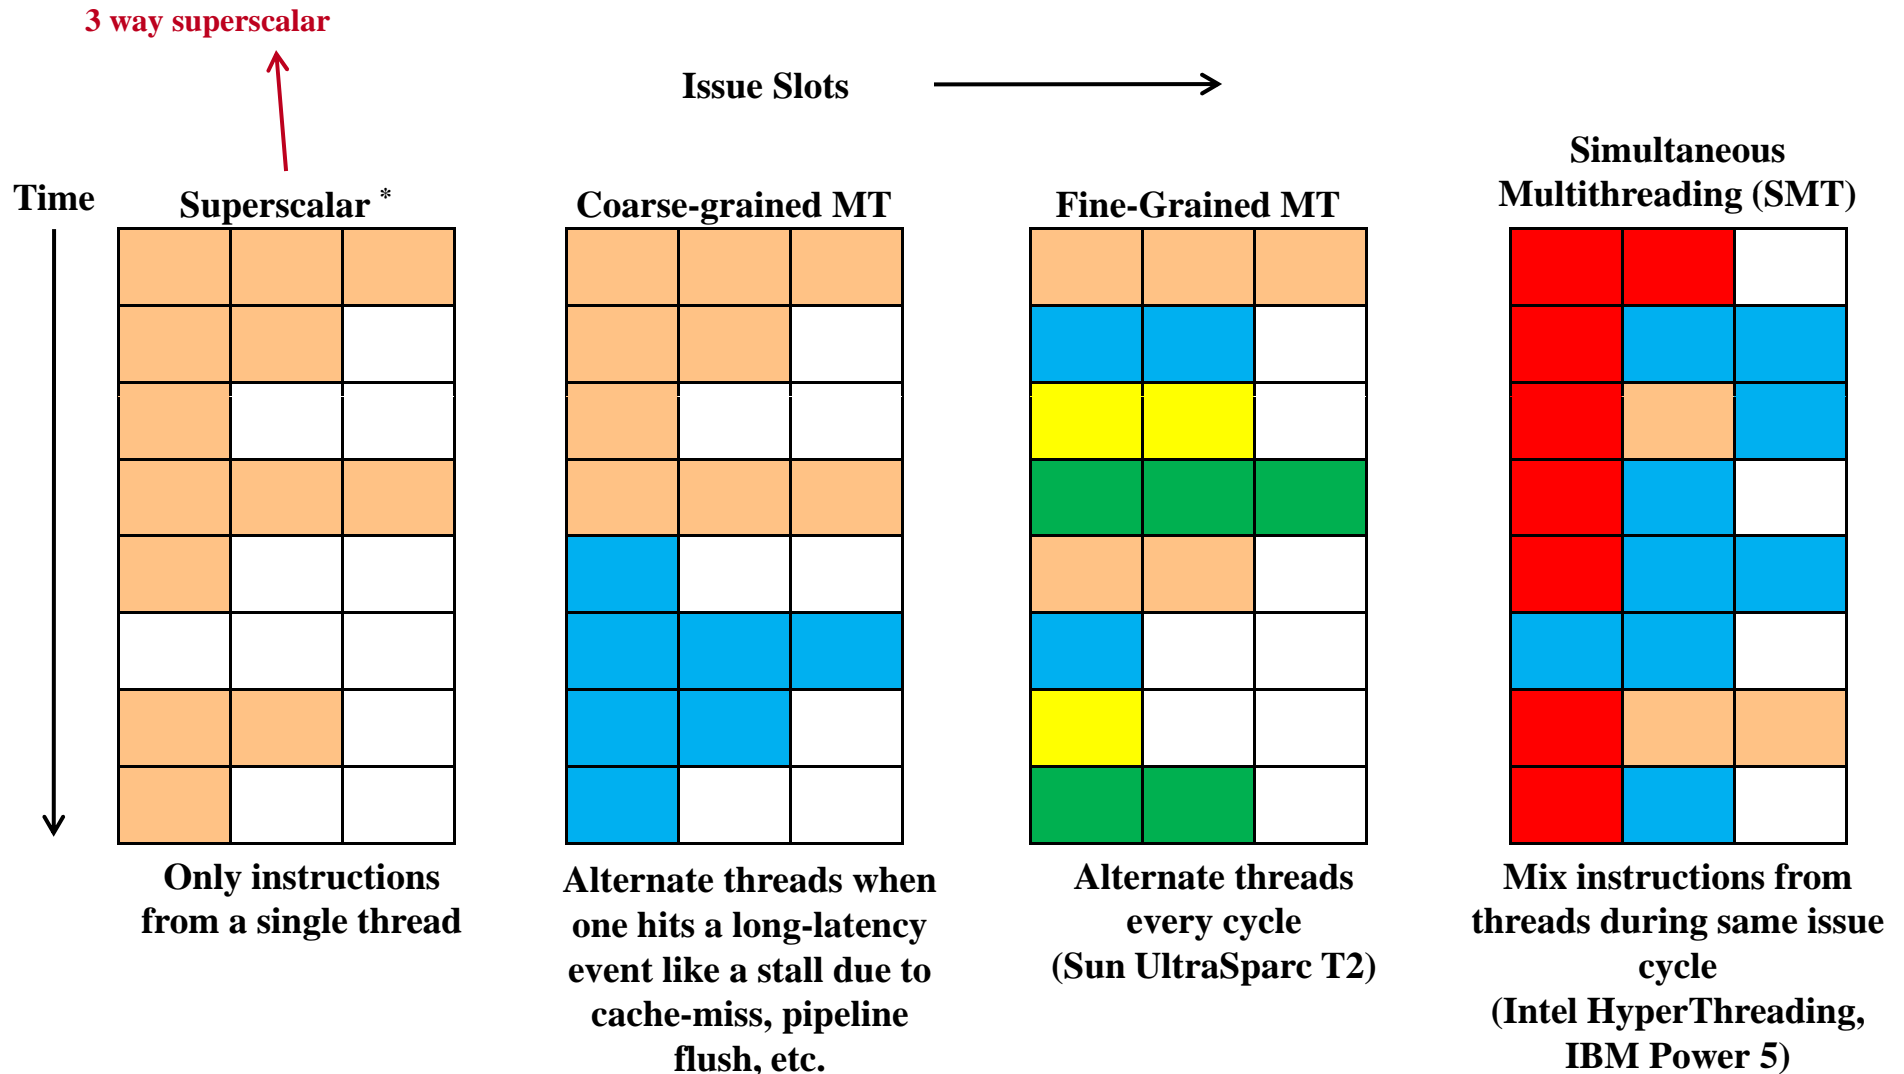
- Instructions are fetched from two different threads
- Different styles of multi-threading determine when they issue (Coarse- or fine-grained, simultaneous)
- There is some level of sharing of resources



Types/Levels of Multithreading

- How should we overlap and share the HW between instructions from different threads
 - **Coarse-grained Multithreading:** Execute one thread with all HW resources until a long latency event, e.g., cache-miss or misprediction will incur a stall or pipeline flush; in such case switch to another thread
 - **Fine-grained Multithreading:** Alternate fetching instructions from a different thread each clock
 - **Simultaneous Multithreading:** Fetch and execute instructions from different threads at the same time


Levels of TLP



*** Discussed later**

Performance Insights

- Parallelism can be exploited in time & space
 - Time = Deep pipelining with HIGH clock frequencies
 - Space = Multiple ways / Multi-threading
 - HW needed for an n-way issue processor requires HW to grow as $O(n^2)$
- Improving performance of a single thread
 - Requires wide issue processor with fast clock rate
- All of this is BAD for power consumption!!



Multithreading Performance

- Takes better advantage of the multiple issue slots and functional units of a modern processor
- Coarse or fine-grained MT
 - Only issue instructions from one thread per clock (limited by ILP)
 - Single thread performance is lowered because of time division between threads
- Simultaneous MT
 - Good single thread performance while filling 'idle' slots with another thread
 - Requires all HW resources of a multiple-issue processor
- However, n-way issue processor requires HW to grow as $O(n^2)$ which is bad for power consumption

Power

- Power consumption decomposed into:
 - **Static:** Power constantly being dissipated (grows with # of transistors)
 - **Dynamic:** Power consumed for switching a bit (1 to 0)
- Most power consumption is dynamic
- $P_{\text{DYN}} = I_{\text{DYN}} * V_{\text{DD}} \approx \frac{1}{2} C_{\text{TOT}} V_{\text{DD}}^2 f$
 - Recall, $I = C \, dV/dt$
 - V_{DD} is the logic '1' voltage, f = clock frequency
- Reducing voltage and frequency reduces power
 - Unfortunately, higher frequencies require higher voltages
 - But, reducing frequency also allows us to reduce V_{DD}
 - Implies power is proportional to f^3 (a cubic savings in power if we can reduce f)

Temperature

- Temperature is related to power consumption
 - Locations on the chip that burn more power will usually run hotter
 - Locations where bits toggle (register file, etc.) often will become quite hot especially if toggling continues for a long period of time
 - Too much heat can destroy a chip
 - Can use sensors to dynamically sense temperature

Temperature (Cont.)


- Techniques for controlling temperature
 - **External measures:** Remove and spread the heat
 - Heat sinks, fans, even liquid cooled machines
 - **Architectural measures**
 - Throttle performance (run at slower frequencies / lower voltages)
 - Global clock gating (pause..turn off the clock)
 - None...results can be catastrophic
- Optional to review:
<http://www.tomshardware.com/reviews/hot-spot,365.html>

MT Performance Insights

- Change goal from pure performance to performance/HW area or performance/watt
- Multithreading is a step in the right direction but can be taken further by replicating more and more processor cores (which may themselves be multithreaded)
- We can increase total computing throughput by using multiple processor cores on a chip
 - Can keep frequency down while still achieving good overall computing throughput

Chip Multiprocessors (CMP)


- Separate cores (duplicating all resources) w/ low issue widths
 - Avoids time division of fine- and coarse-grained MT
 - Often requires simpler, smaller HW than MT at similar issue slots
 - A 16-way issue superscalar/SMT would likely take more area and be MUCH harder to design than 16 simple, separate cores
 - Better
 - performance / area
 - performance / watt
- Greater “spatial” parallelism (via the many cores) allows us to reduce clock rates (reduce temporal parallelism)




Intel® Core 2 Duo™ Processor Die Photo

Typical CMP Organization

- Processors each have their own L1 caches (always access data from these caches)
- Shared L2 or memory
- Interconnection network for communicating data between caches
 - Shared Bus allows only one transfer at a time
 - Switched interconnection network allows multiple simultaneous transfers



**Simple CMP connected via shared bus
(Core 2 Duo)**



**More cores usually requires a more
advanced interconnection network
(Cell Processor)**