

9/27/10

CPU Scheduling

Task: Which thread gets the CPU task - & when do they get it

Policy: Rules }

Mechanism: ~~Implementation~~ : How we implement the rules }

Metaphor

Policy: Traffic laws

Mechanism: Police

5 Possible Goals for Schedulers

Fairness: Every job gets ^{the same} amount of CPU time

Efficiency: Keep as many ~~as~~ of the computer's resources as busy as possible

Response Time: Minimize / Maximize interactive response time

Turnaround Time: Give CPU preference (or not) to background tasks

Throughput: Complete as many jobs as possible in some amount of time

Issue: Do jobs run to completion,
or not.

└ non-preemptive

└ preemptive

✓

How long is a time slice?

too long: interactive use will
suffer

too short: too many unnecessary
context switches

Scheduling Policies

First-Come First-Served

- Similar to "standing in line"
- Non-preemptive

- The job that has waited the longest gets the CPU next

+ Fair (No starvation)

+ Simple

- Short jobs can get "stuck" behind long jobs

Round Robin

Preemptive FCFS

Jobs run until their time slice has expired, ~~or~~ they are blocked on a "slow operation", or they voluntarily give up the CPU.

+ Fair

+ Short jobs don't get stuck

behind long jobs

- Unnecessary context switches when all jobs are about the same length (not short)

Shortest Job First

Maximizes throughput

Run the job requiring the
least ~~complet~~ CPU time to
complete }

Non-preemptive.

— Not Fair (Starvation)

— Not implementable

Shortest Remaining Time to Complete First

- Preemptive

The job "we" predict ~~requiring~~ will use the smallest time slice percentage - based on its ~~past~~ recent history

- Can still have starvation

- The scheduler must track & compute the time slice usage

Priority-Based Scheduling

Priority: Some jobs are favored over other jobs

Policy: "Higher" priority jobs run before lower priority jobs

Issue: Preemptive or not?

New Issue: Do priorities change over time?

Static priorities: Priorities are fixed

- Starvation
- + Easier on O.S.

Dynamic priorities: Priorities can change

- up or down
- wait time
- CPU time

- + No starvation
- Extra overhead, during context switch, in recomputing priorities

Final Issue: How many priorities are needed?

Somewhere around 4 priorities are typical

Most effective Ready Queue organization, is to have a separate Ready Queue for each priority

- This allows for different scheduling ~~for~~ policies for each priority
- The time slice interval can be different for each priority queue

Deadlock

Occurs because of competition for resources

- hardware
- files
- synchronization primitives

Ex 1

lock1 → A
{
lock2 → A
lock1 → R

{
 lock1 → Acquire();
 lock2 → Acquire();
 cv2 → Wait(lock2);
}

lock1 → Acquire
lock2 → Acquire
cv2 → Signal(lock2)

Ex 2

{
 lock1 → Acquire()
 → {
 lock2 → Acquire()
 }
}

{
 lock2 → Acquire()
 lock1 → Acquire()
}

Define Deadlock

Two, or more, jobs each waiting on an event that can only be produced by one of the waiting jobs.

Using Resource Sequence

- ① User program requests resource access from the O.S.
- ② Once access is granted, the user's job uses the resource
- ③ When user program is done with the resource, it is given back to the O.S.

What if a requested resource
is not available?

① Fail the request

Application has the responsibility
to handle failed requests
+ No deadlock possible

② Queue the request

O.S. has responsibility to manage

queued requests

- Deadlock is possible

4 Conditions For Deadlock

1. Mutual Exclusion

Some resources that cannot be shared

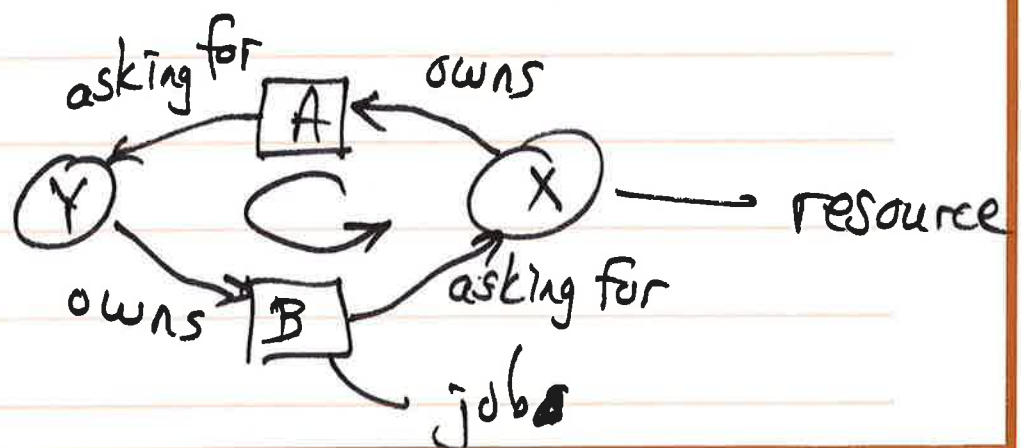
2. Hold & Wait

Jobs can own resources & request other resources, without giving up what they already own.

3. No preemption

Once a resource is given to a job, it cannot be "safely" borrowed

4. Circular Wait



Deadlock Solution Strategies (3½)

1. Do nothing

2. Detection & Recovery

have a way
to detect
deadlock has
occurred

kill a job

3. Dynamic Avoidance

We try to keep deadlock from happening by carefully allocating resources.

4. Prevention

Eliminate 1 of 4 required conditions
for deadlock