

10/4/10

How to approach a system call

Don't start thinking about writing code.

DO think about the purpose of a system call

• "Safe" mechanism by which user programs are able to access kernel-protected resources

Lock class & Condition are kernel resources

Lock & Condition system calls are just interfaces to Lock & Condition kernel objects

## Example: Acquire

Task: Provide access to a Lock object

\* Locks are specific to \*  
a process

The Kernel must have a  
real Lock object - created  
upon request by a user

Program. - CreateLock syscall

Locks must be retrievable  
for Acquire, Release, Destroy Lock  
syscalls

Condition variables require a  
lock object to be passed  
in as an argument

## Format of syscalls from user program viewpoint

What this means:

- ① What data does the OS need to carry out the request
- ② Is there a return value to the user program?

<u>int</u>	CreateLock	} Making a resource for the user program for <u>later use</u>
<u>int</u>	CreateCondition	

• Must return a "handle" to the user program - a resource identifier

Let's use an array (kernel) to store created Locks & CVs

- 1<sup>array</sup> for each - 1 for Locks  
1 for CVs

My kernel return value is the index position • in the array of the new object



Because of my choice of an array,  
I have to track of where the  
next index position is for the  
next new resource

- How do I know when I run out of space in my array

\*

- Use a simple counter.

- Every time a Create syscall is requested I increment

the counter

- Be sure to check for a full array

The tables & counters are shared  
by all threads issuing system  
calls

- Need a <sup>kernel</sup> lock - used by the kernel - to protect these shared kernel resources

The array cannot be an array of Locks/Conditions

- We cannot enforce the process-specific rule
- Need some piece of data in the array so the kernel can enforce the process rule
  - If you have a unique PID in your process

table you can use it

- Or... AddrSpace pointer

For now...

```
struct LockEntry {  
    Lock *lock;  
    AddrSpace *ownerProcess;  
};
```

```
LockEntry lockTable[MAX-LOCKS];
```

One last Key point: DestroyLock

Task: Delete an existing Lock

Issue: One thread, in a process,  
issues a DestroyLock ~~before~~  
~~the~~ while the Lock is  
in use

OS Help: Only destroy a resource

when it is not in use.

What does it mean for a Lock to be in use?

- state should be "busy"

Rule: IF I <sup>OS</sup> get a DestroyLock request,  
I only delete the lock when  
the lock state is "available"

os  
A new issue: If we don't "remember"  
a Destroy Lock has been  
requested, the user  
program would be forced  
to issue - worst case -  
multiple Destroy Locks  
to delete a lock.

new

```
struct LockEntry {  
    Lock# lock;
```

```
    AddrSpace* ownerProcess;  
    bool toBeDeleted = false;  
};
```



## Signature of all Lock syscalls

int Create Lock (char \*, int len)  
void Acquire (int)  
void Release (int)  
void Destroy Lock (int)

index into the kernel lock table for created lock

char \*, int len

```
int Create Condition (char *, int len)
void Wait (int cond#, int lock#)
" Signal (int cond#, int lock#)
" Broadcast (int cond#, int lock#)
void Destroy Condition (int cond#);
```



syscall.h has:

\* void Halt();      function declaration

To the compiler: There is a function called Halt

In my user program, I can  
#include "syscall.h"

```
{  
int main() {  
    Halt();  
    return 0;  
}
```

① compile  $\Rightarrow$  object file

② Link  $\Rightarrow$  links all object files w/ any other libraries to make an executable

There is C/C++ function  
called `Halt()` - it is just a  
"stub"

The connection is in `start.s`.  
There is assembler code that  
"converts" the stub function  
into a jump to Nachos kernel

# The pattern for syscall implementation

① Get the user program arguments  
int \_\_\_\_ = machine → ReadRegister(4)  
5  
6  
7

② Validate the arguments

If any argument is invalid,  
print an appropriate error

msg and return

③ Perform the task

④ If called for, return a value

## Acquire

lockTableLock → Acquire();

• validations

Lock\* lock = lockTable[Index] → lock;

~~isToBeUsed~~ = true;

(isInUse)

Do the  
task

lock → Acquire();

lockTableLock → Release();

Need a status variable that a  
lock is going to be used



Scenario: A lock is not in use  
BUT the thread next  
in the Ready Q has  
an Acquire about to  
be executed

The thread in the CPU has  
requested a Destroy Lock  
• ~~before~~ after validation  
this thread is context

switched out

This cannot be allowed to happen  
Destroy Lock

lockTable Lock → Acquire();  
- validation

context  
switch



• "delete" the entry

lockTable Lock → Release();